



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2019

Anomaly Detection for Cloud Native Storage

Xing Yang, OpenSDS and Seiya Takei, Yahoo! JAPAN

Agenda

- Introduction to anomaly detection
 - Anomaly detection in storage performance
- Yahoo! JAPAN use cases
- Telemetry
 - Integrate with Prometheus and Grafana
 - Collect performance metrics from storage backends
 - Metrics drivers: LVM, Ceph, other storage systems...
- Anomaly detection
 - Detect anomalous data points based on metrics collected from Telemetry.
- Demo

What is Anomaly Detection

- Anomaly detection is a technique used to identify unusual patterns that do not conform to expected behavior, called outliers.
- Categories of anomalies:
 - Point anomalies
 - Contextual anomalies
 - Collective anomalies



Anomaly Detection Use Cases

- Intrusion detection
 - identifying strange patterns in network traffic that could signal a hack
- Medical health
 - spotting a malignant tumor in an MRI scan
- Fraud detection
 - credit card, cell phone, insurance claim fraud, etc.
- Fault detection
 - mechanical units, etc.
- Anomaly detection in storage
 - disk failure, etc.



Storage Performance Challenges



KubeCon

CloudNativeCon



CloudNativeCon

China 2019

Bottlenecks

- Disk failure/Inaccessible disks
- Read/Write I/O errors
- Volume issues
- Port masking
- Configuration issues – Host, Storage subsystem, port, Interoperability
- Network congestion
- Workload configurations
- UPS battery failure
- Port protocol errors,
- Port congestion

Metrics

- I/O Rate R/W,
- Data Rate R/W,
- Response time R/W,
- Cache hit R/W,
- Data block size R/W,
- Porta data rate R/W,
- Port-local node queue time

Correlations

- CPU & Network Traffic
- CPU & Memory
- Port & Host counters
- IOPs, read rate, & CPU, memory

Source: [Using Machine Learning for Intelligent Storage Performance Anomaly Detection](#)



Yahoo! JAPAN's Environment

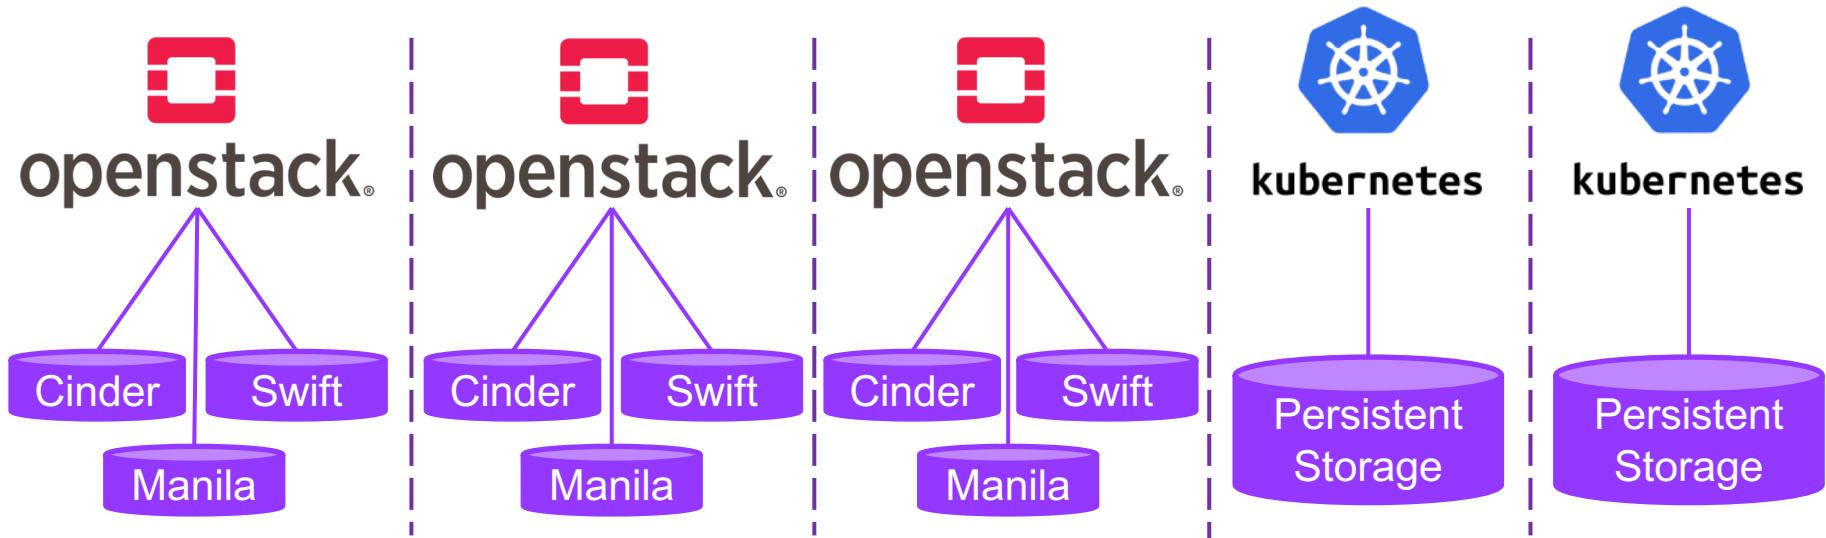
- Private Cloud
 - IaaS
 - 140,000 VMs
 - PaaS
 - 30,000 Containers
 - CaaS
 - 390 Kubernetes clusters



kubernetes

Storage for Private Cloud

- OpenStack/Kubernetes clusters each have storage
- There are many storages in Yahoo! JAPAN's environment



Future environment we hope

- We want to manage storages using OpenSDS
- If telemetry and anomaly detection can be managed by OpenSDS, we can manage many storages easier



openstack.



openstack.



openstack.



kubernetes



kubernetes

OpenSDS



Software Defined Storage

- OpenStack/kubernetes backend storage
- Running SDS in Yahoo! JAPAN's environment
 - Ceph
 - Using in test environment
 - Quobyte
 - Started using it end of last year
- Difficult to operate
 - Health management of distributed system
 - server, network, ...
 - Telemetry and Anomaly Detection are very important



The Open Autonomous Data Platform

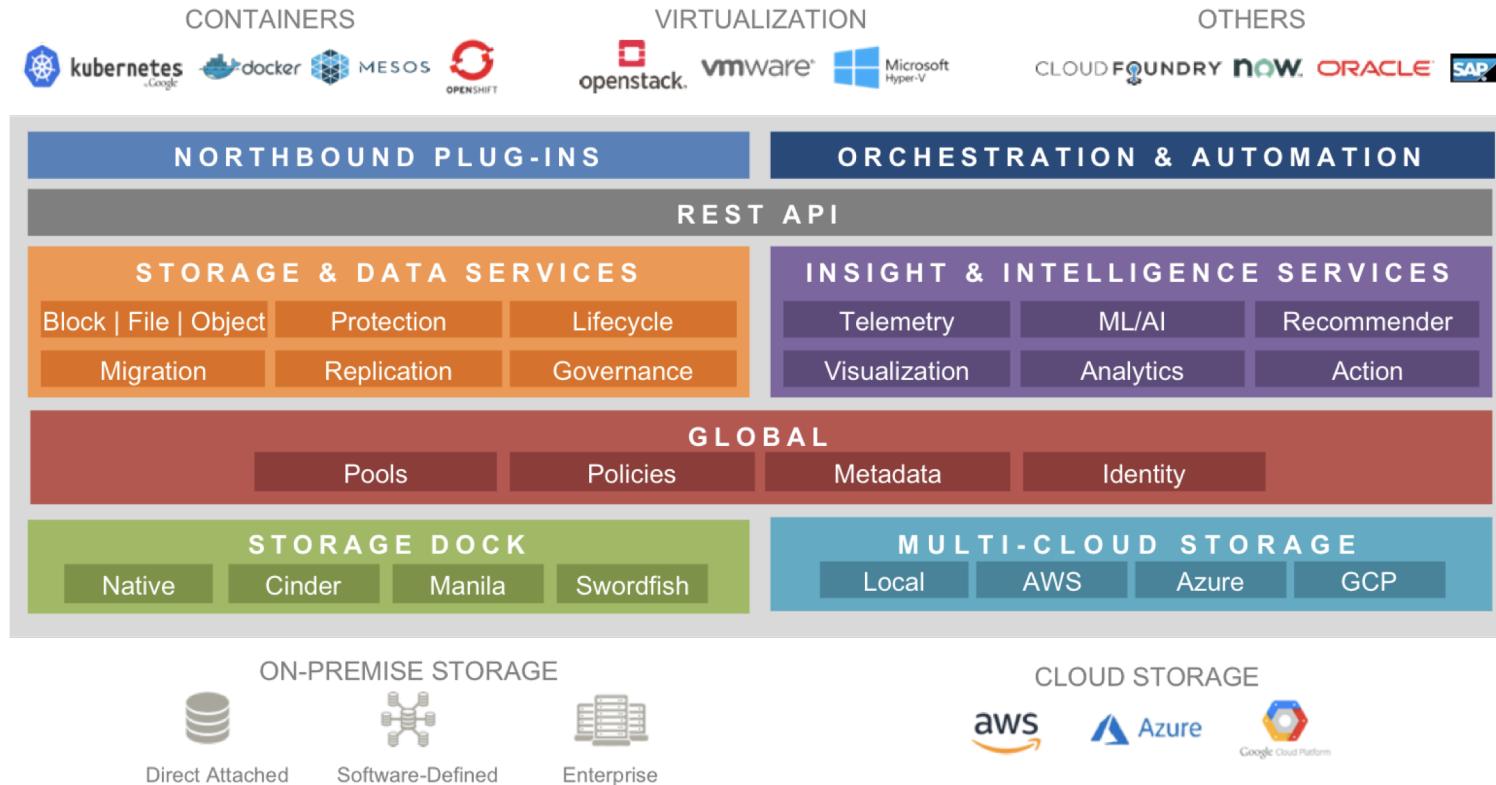


KubeCon

CloudNativeCon

OPEN SOURCE SUMMIT

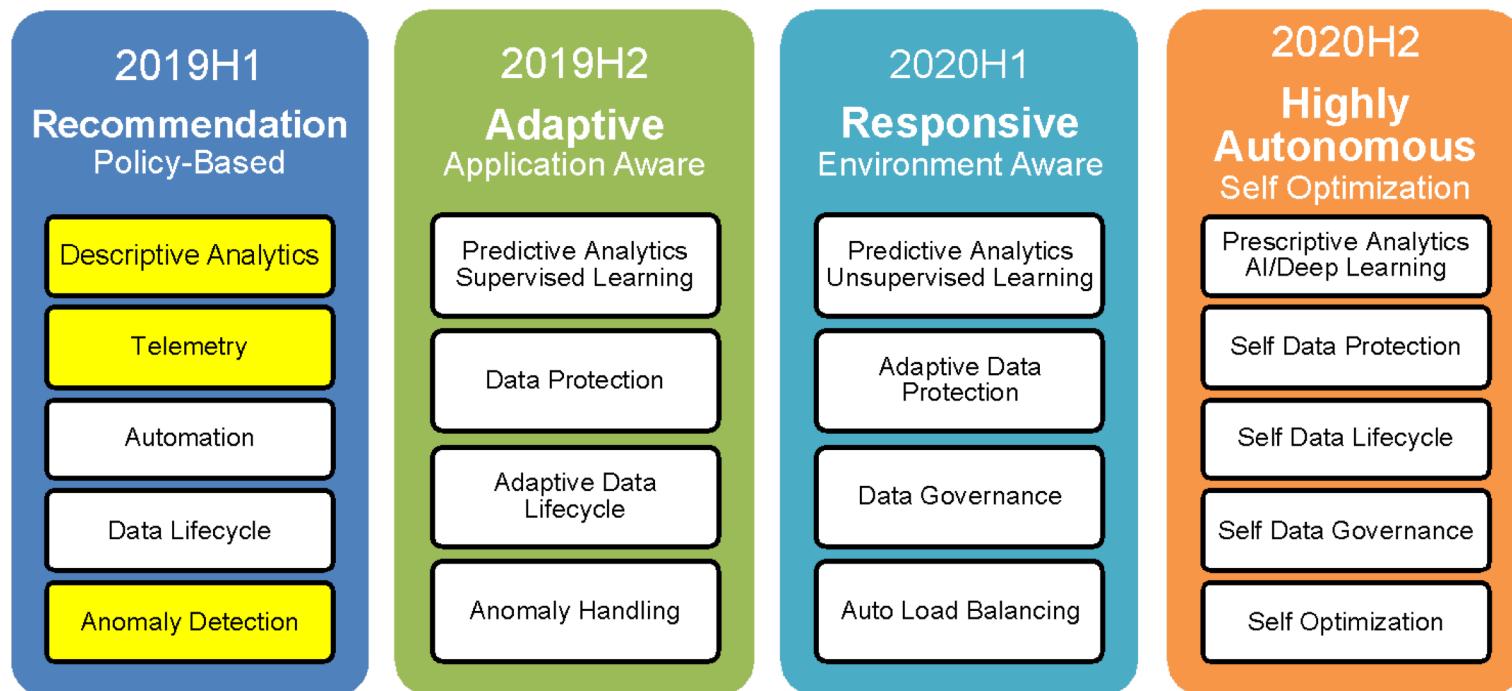
China 2019

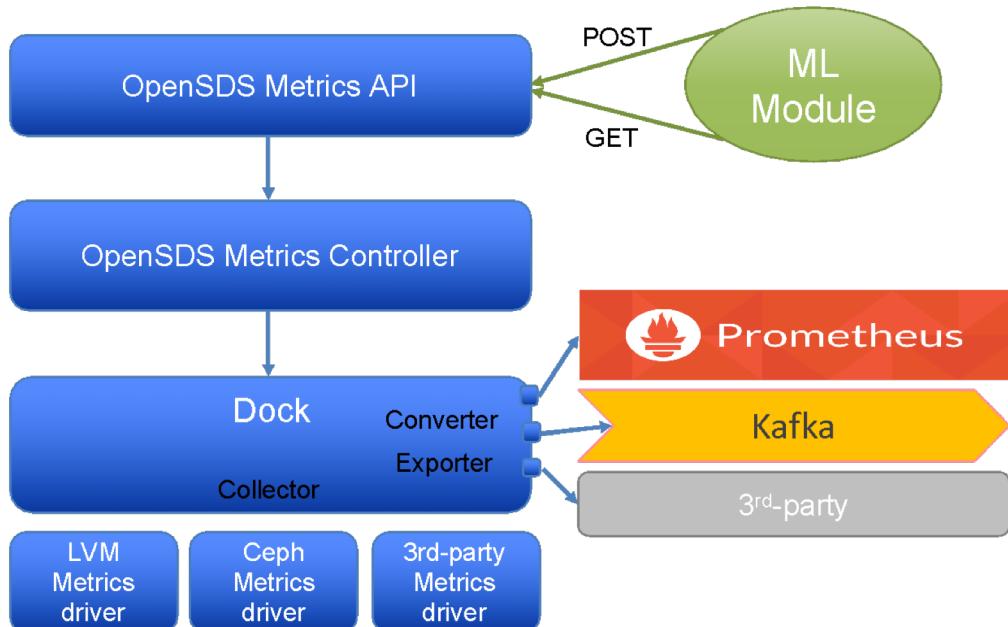


The Road To Autonomous Data Platform



CloudNativeCon
OPEN SOURCE SUMMIT
China 2019





- ML module sends requests using Metrics API that generates data.
- Collector collects metrics from metrics drivers.
- Adapter includes a Converter that converts data to a proper format that can be understood by the receiving end, e.g., Prometheus, and an Exporter that sends(emits) the data to the intended destination.
- ML module receives data through Kafka. ML module also retrieves additional data using Metrics API which gets data from Prometheus.
- Collected metrics include IOPs, bandwidth, latency, average CPU usage, etc. for various resources such as storage controller, pools, volumes, disks, etc. For Ceph, an existing Prometheus Ceph exporter will be used. Prometheus Node exporter will also be used to collect node metrics.
- OpenSDS dashboard is integrated with Grafana to display metrics and Prometheus Alert Manager to show alerts.

Collecting Storage Performance Metrics

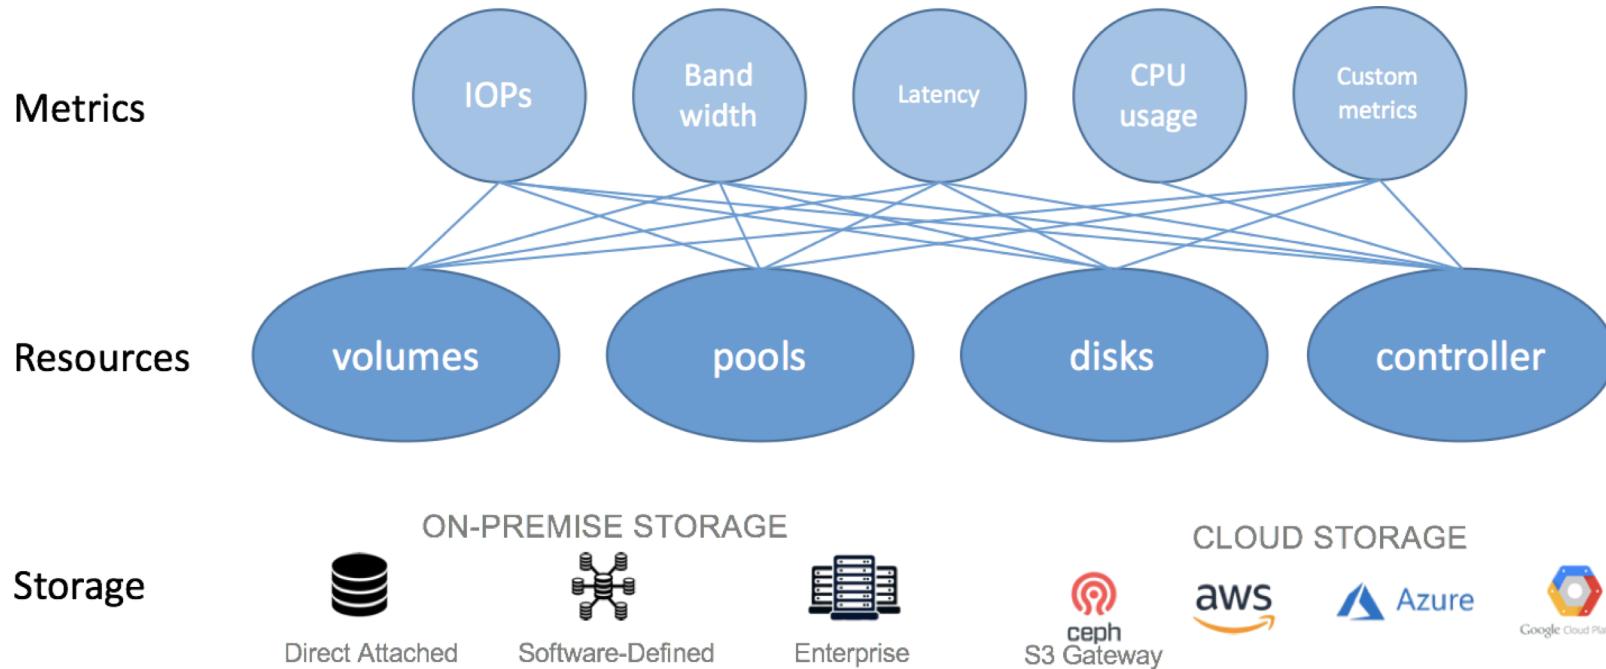


KubeCon

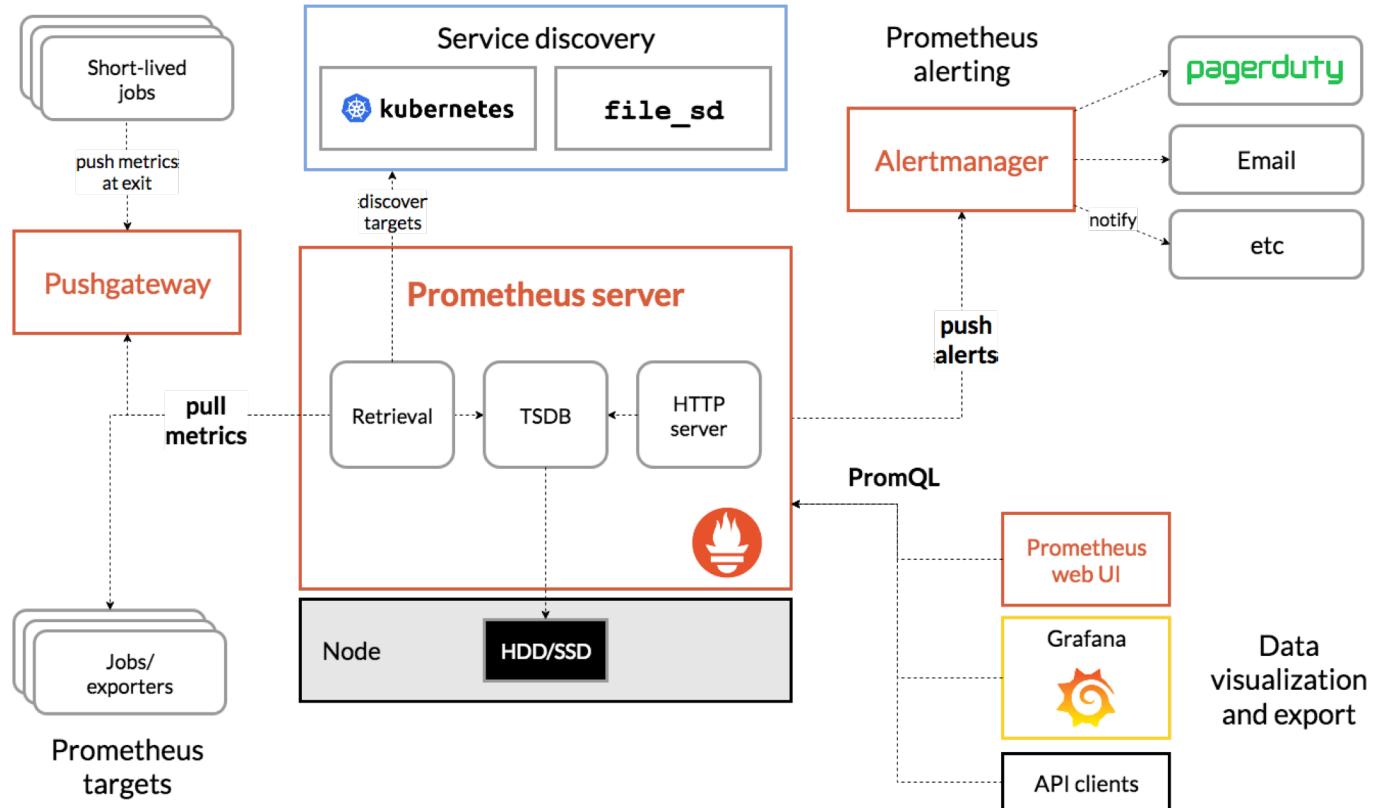
CloudNativeCon

OPEN SOURCE SUMMIT

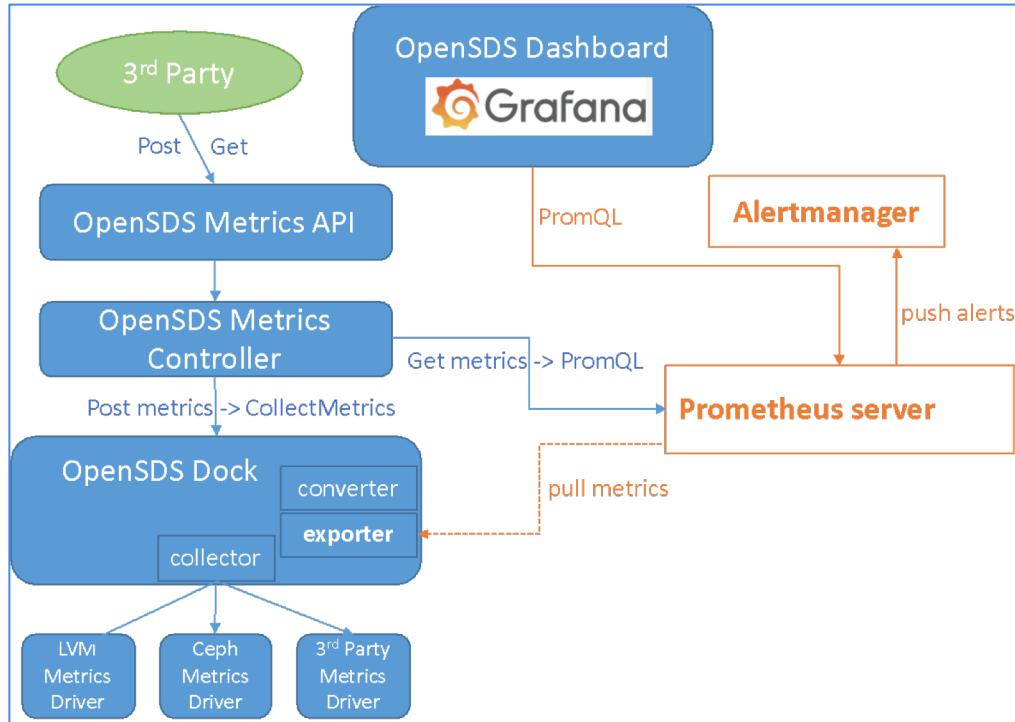
China 2019



Prometheus Architecture



Emit Metrics to Prometheus



- Post request will be sent to the Metrics driver to collect metrics
- Get request will be re-routed to Prometheus server using PromQL
- Metrics will be saved in Prometheus database.

Metrics interface

```
func CollectMetrics() ([]*model.MetricSpec, error)
```

```
type CollectMetricSpec struct {  
    *BaseModel  
  
    DriverType string  
}
```

```
type GetMetricSpec struct {  
    *BaseModel  
  
    InstanceId string  
    MetricName string  
    StartTime string  
    EndTime string  
}
```

```
type MetricSpec struct {  
    InstanceID string  
    InstanceName string  
    Job string  
    Labels map[string]string  
    Component string  
    Name string  
    Unit string  
    AggrType string  
    MetricValues []Metric  
}  
  
type Metric struct {  
    Timestamp int64  
    Value float64  
}
```



Collect LVM Metrics



KubeCon
OPEN SOURCE SUMMIT



CloudNativeCon
China 2019

- Tools:
 - lvmsar (LVM system activity reporter)
 - iostat

- resource: volume
- metrics:
 - iops (tps)
 - read_throughput (kb/s)
 - write_throughput (kb/s)
 - response_time (ms)
 - service_time (ms)
 - utilization_percentage (%)

- resource: disk
- metrics:
 - iops (tps)
 - read_throughput (kb/s)
 - write_throughput (kb/s)
 - response_time (ms)
 - service_time (ms)
 - utilization_percentage (%)

Collect Ceph Metrics

- Use existing Ceph exporter in Prometheus

 <https://prometheus.io/docs/instrumenting/exporters/>

- RabbitMQ Management Plugin exporter

Storage

- Ceph exporter
- Ceph RADOSGW exporter
- Gluster exporter
- Hadoop HDFS FSImage exporter
- Lustre exporter
- ScaleIO exporter

HTTP

- Apache exporter
- HAProxy exporter (official)
- Nginx metric library
- Nginx VTS exporter

- resource: pool
- metrics:
 - pool_used_bytes
 - pool_available_bytes
 - pool_objects_total
 - pool_dirty_objects_total
 - pool_read_total
 - pool_read_bytes_total
 - pool_write_total
 - pool_write_bytes_total

- resource: cluster
- metrics:
 -

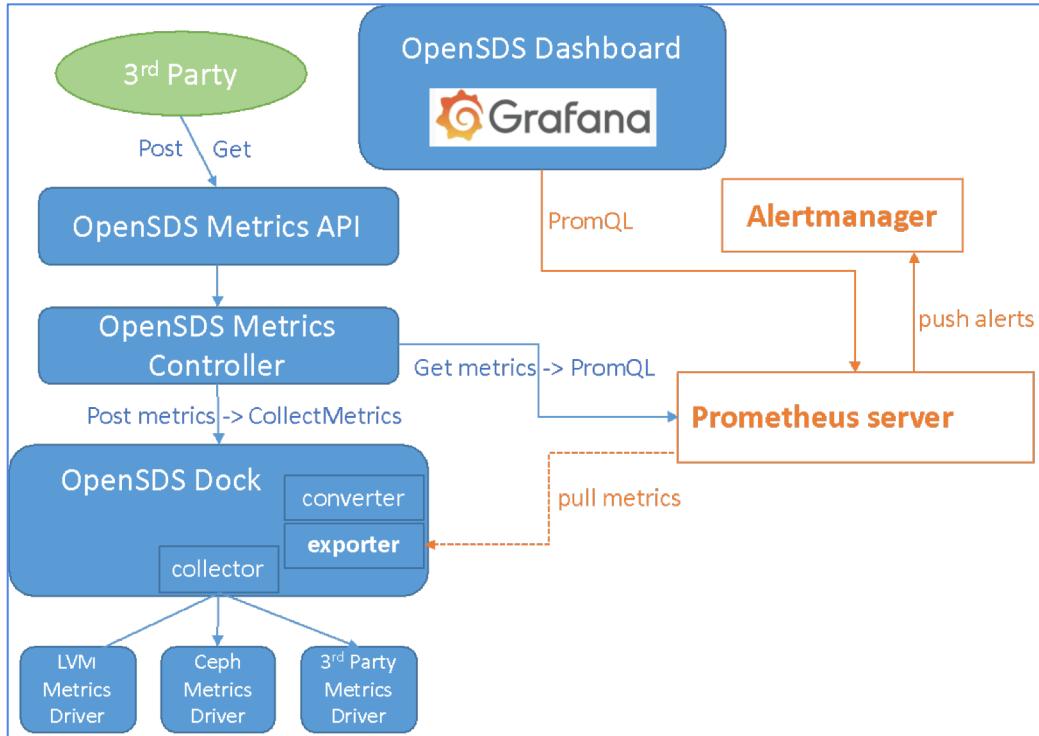
LVM Metrics in Grafana



LVM Metrics in Grafana



Emit Metrics to Prometheus (recap)



- Post request will be sent to the Metrics driver to collect metrics
- Get request will be re-routed to Prometheus server using PromQL
- Metrics will be saved in Prometheus database.

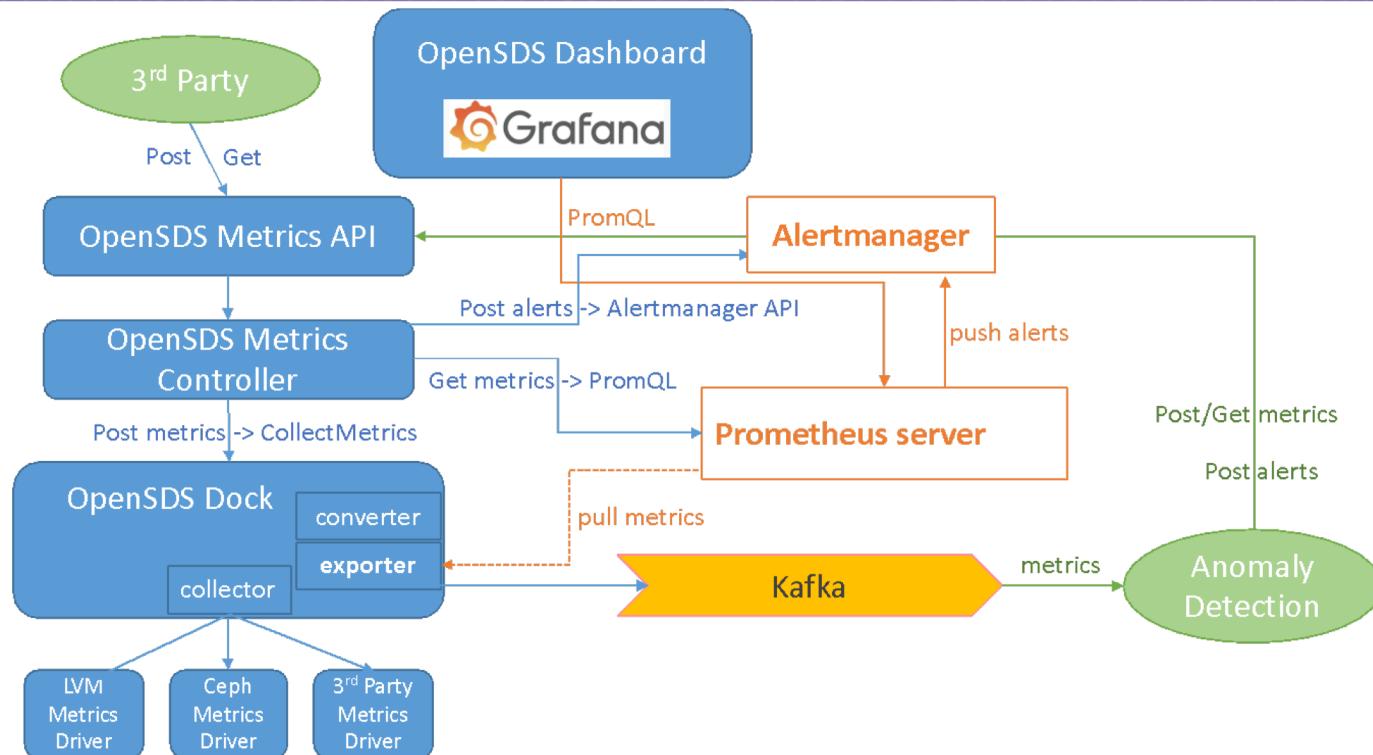
Prometheus Alert Manager

- The alert rules are configured in Prometheus, and on the thresholds being crossed, Prometheus will raise an alert to the Alertmanager. Alerts can be defined on raw metrics and derived metrics.
- The Anomaly Detection module detects an anomaly, raises a custom alert to Alertmanager using the REST API interface of Alertmanager

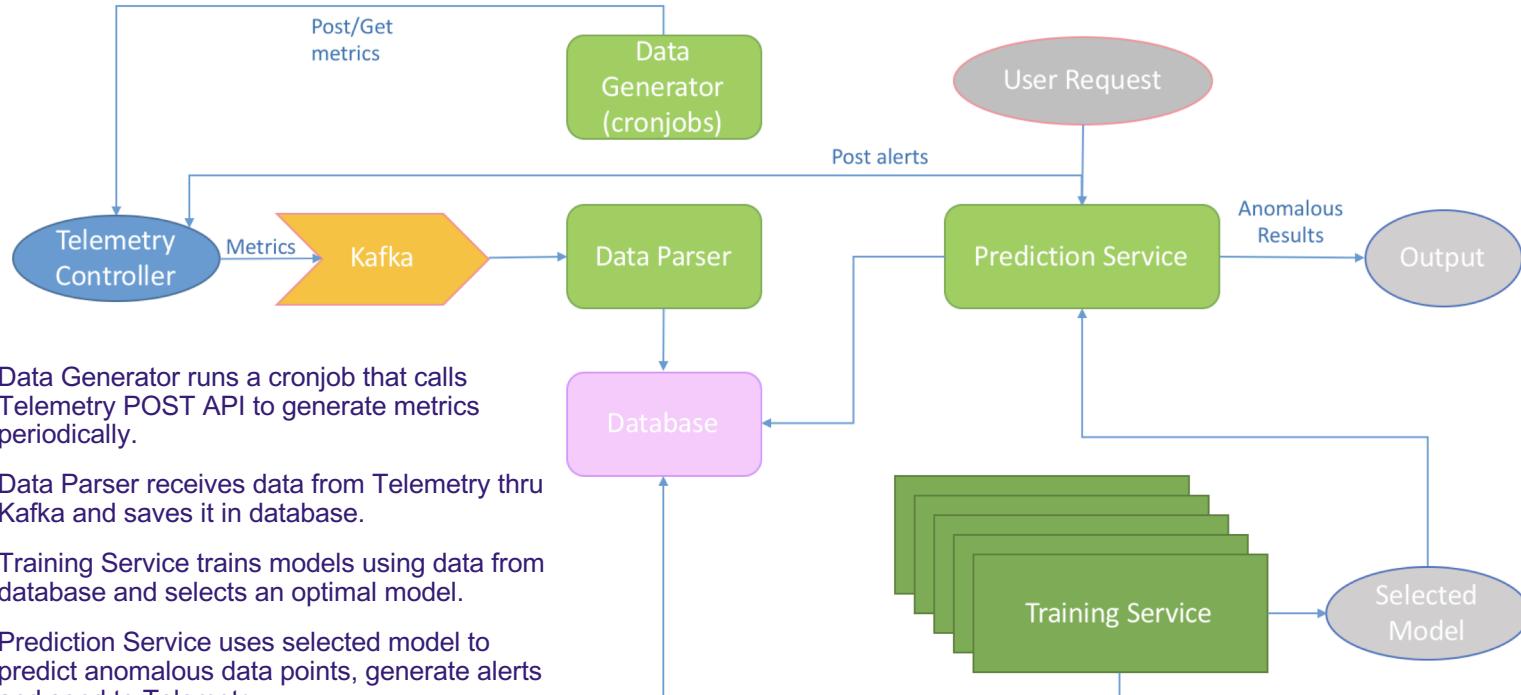
The screenshot shows the Alertmanager interface with two alerts listed:

- Alert 1:** alertname="DiskRunningFull" (14:34:23, 2019-03-28 UTC)
Custom matcher: env="production"
Labels: dev="sda1", instance="example1"
Annotations: + Info, Silence
- Alert 2:** alertname="DiskSpace10%Free" (14:23:34, 2019-03-28 UTC)
Custom matcher: env="production"
Labels: device="/dev/xvda5", fstype="ext4", instance="localhost9100", job="node_exporter", mountpoint="/", severity="moderate"
Annotations: + Info, Source, Silence

Send Metrics through Kafka



Anomaly Detection Architecture



Anomaly Detection Algorithms



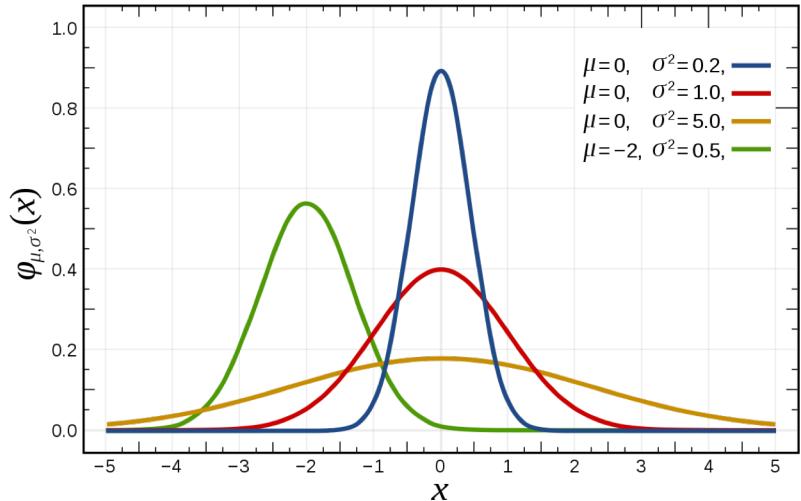
- Classification based
 - A classifier that can distinguish between normal and anomalous classes can be learned in the given feature space.
- Nearest neighbor based
 - Normal data instances occur in dense neighborhoods, while anomalies occur far from their closest neighbors.
- Clustering based
 - Normal data instances belong to a cluster in the data, while anomalies either do not belong to any cluster.
- Information theoretic
 - Anomalies in data induce irregularities in the information content of the data set.
- Spectral
 - Data can be embedded into a lower dimensional subspace in which normal instances and anomalies appear significantly different.
- Statistical models
 - **Gaussian model**
 - Regression model



Reference: [Anomaly Detection : A Survey](#)

Gaussian Model

- The Gaussian model is a statistical model that assumes the pattern of the dataset follows the gaussian distribution.
- A threshold needs to be specified to differentiate between normal and abnormal data points.



Source: [Normal Distribution](#)

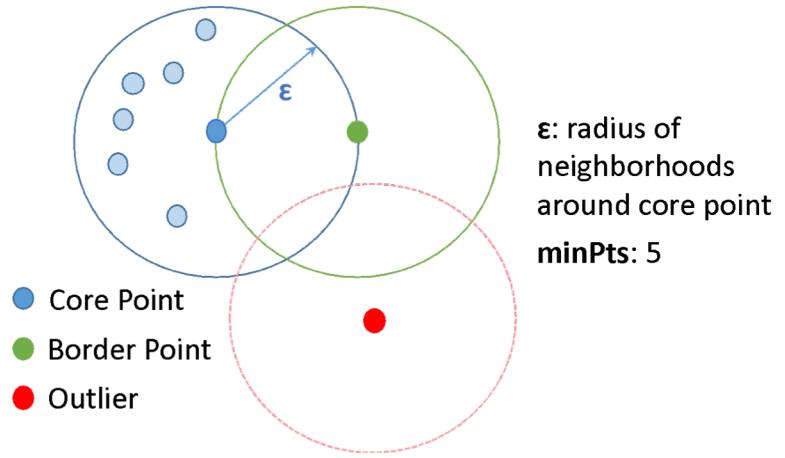


Anomaly detection

Anomaly detection using the multivariate Gaussian distribution by Andrew Ng

DBSCAN Clustering

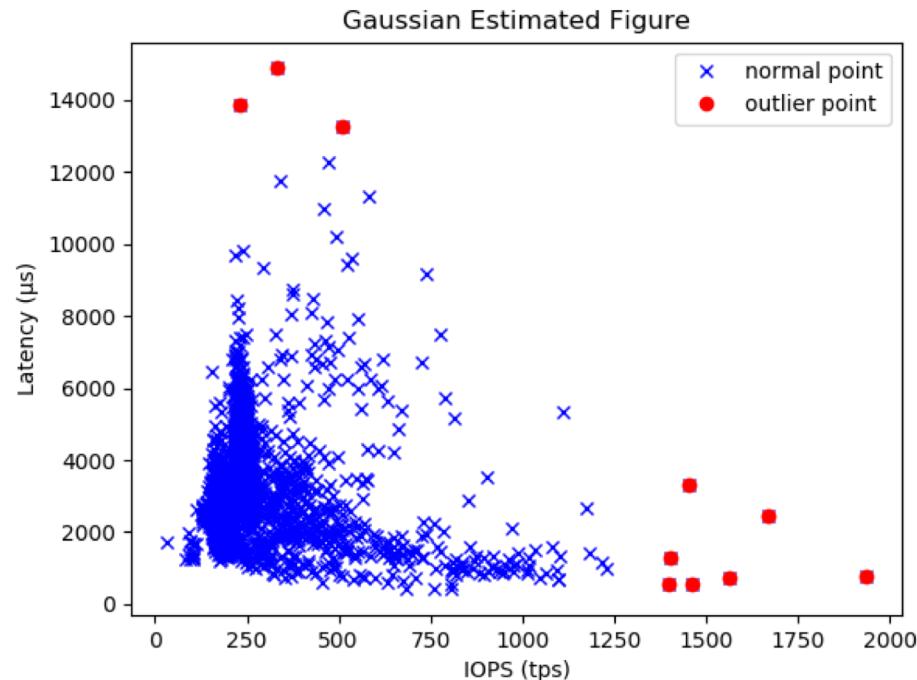
- DBSCAN refers to Density-Based Spatial Clustering Applications with Noise. Clustering is used to group similar data instances into clusters. DBSCAN is a clustering model designed to discover clusters of arbitrary shape based on density.
- The algorithm has two input parameters ϵ and **minPts**.



DBSCAN categories the data points into three categories: Core Points, Border Points, and Outlier.

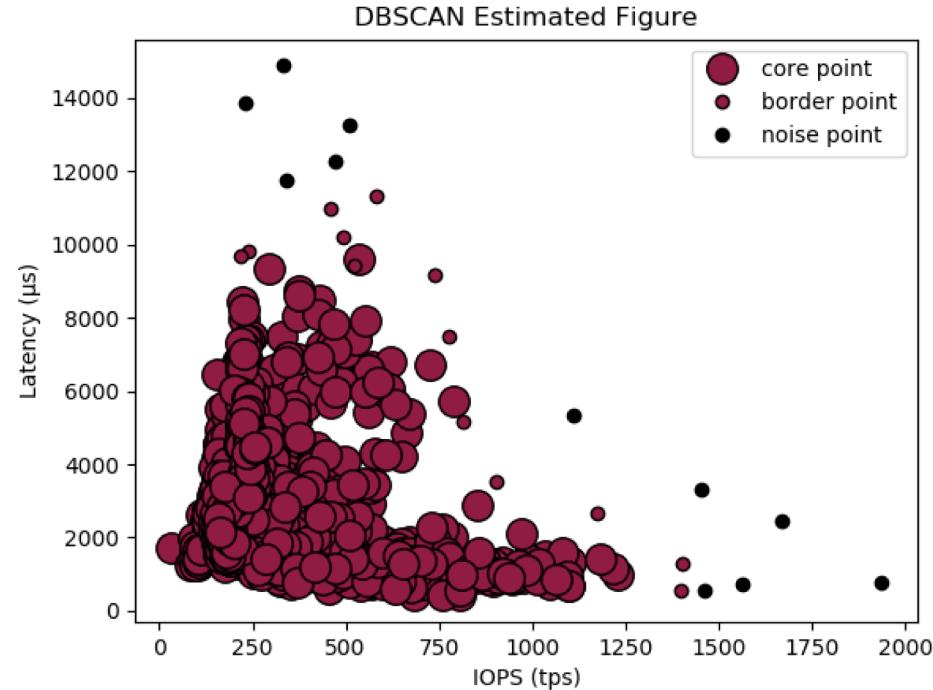
Gaussian Model Graph – LVM

- Volume level metrics
 - IOPs
 - Latency
- Data generation: used dd for random writes and bonnie++ for heavy workloads
- Data processing: 6000 data points collected in 2 days, removing zeros.
- mean: [259.76 3105.44]
- covariance: -16620.57
- epsilon: -41.80
- f1_score: 0.8



DBSCAN Graph – LVM

- Volume level metrics
 - IOPs
 - Latency
- Data generation: used dd for random writes and bonnie++ for heavy workloads
- Data processing: 6000 data points collected in 2 days, removing zeros
- epsilon: 1.40
- minPts: 11
- adjusted_rand_score: 0.69





POC Setup

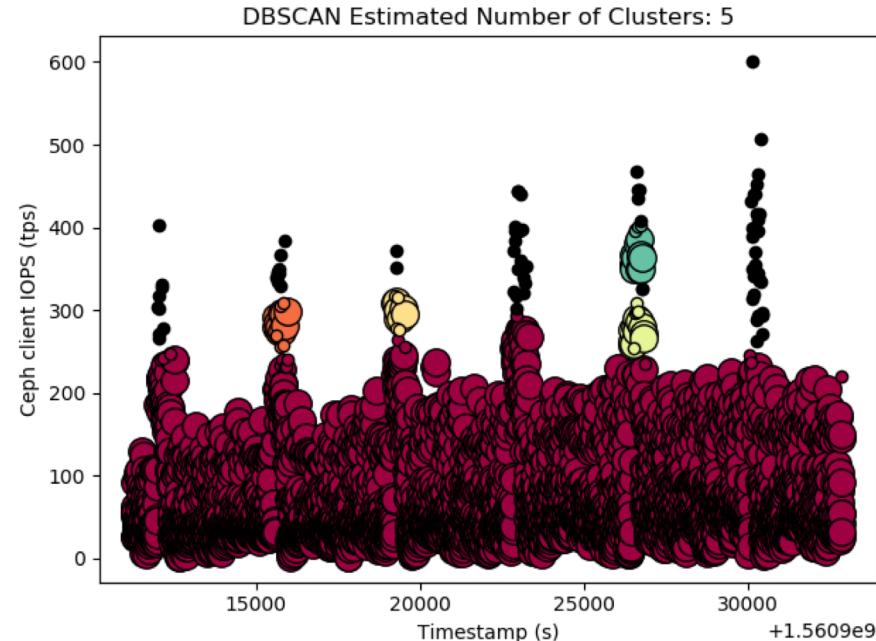
- MongoDB running on a volume provisioned by OpenSDS CSI plugin in Kubernetes environment with Ceph backend
- Collect metrics from Ceph, node-exporter, and MongoDB

```
.....  
volumeClaimTemplates:  
  - metadata:  
      name: mongodb-persistent-storage-claim  
    annotations:  
      volume.beta.kubernetes.io/storage-class: "csi-sc-opensdsplugin"  
    spec:  
      accessModes: [ "ReadWriteOnce" ]  
      resources:  
        requests:  
          storage: 1Gi  
.....
```

```
.....  
apiVersion: apps/v1beta1  
kind: StatefulSet  
metadata:  
  name: mongod  
spec:  
  serviceName: mongodb-service  
  replicas: 3  
  template:  
    metadata:  
      labels:  
        role: mongo  
        environment: test  
      replicaset: MainRepSet  
      name: mongo  
    spec:  
.....
```

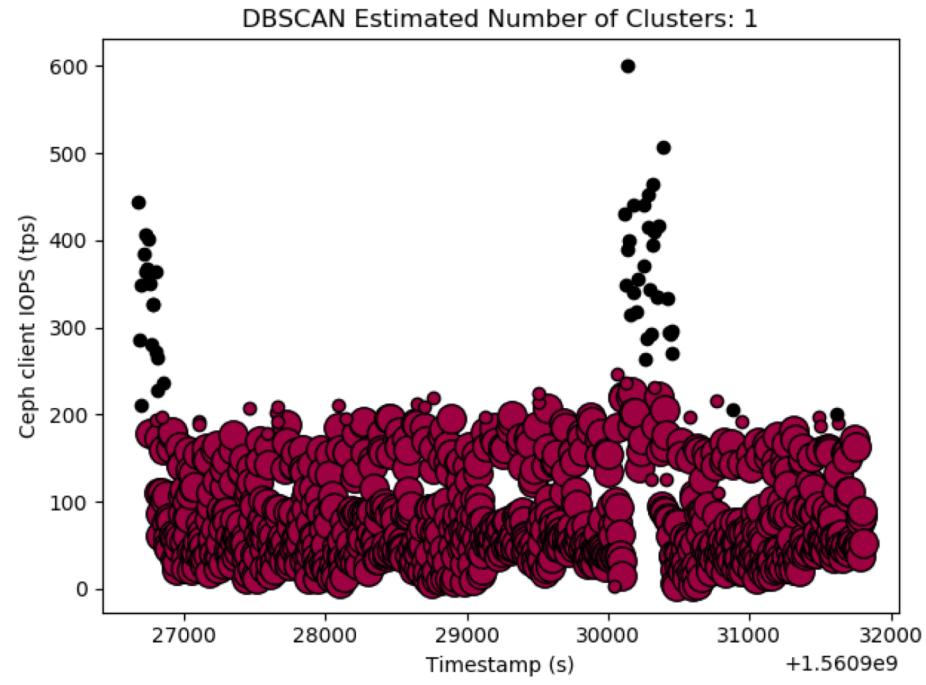
DBSCAN Graph - Ceph

- Ceph Client IOPs
- Data generation: MongoDB tools [Workload Driver](#) and [MongoDB Multithreaded Performance Test Tool](#) were used to generate MongoDB workloads (insert/delete/update/find)
- Data processing: 6000 data points collected in 6 hours, removing zeros
- epsilon: 0.30
- minPts: 10
- adjusted_rand_score: 0.77



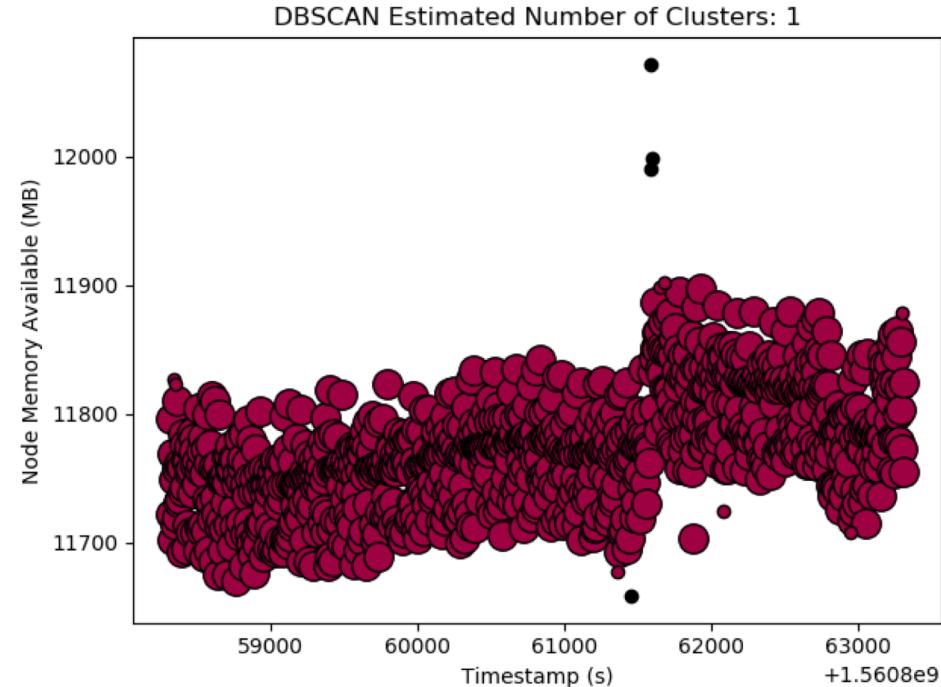
DBSCAN Graph - Ceph

- Ceph Client IOPs
- Data generation: MongoDB tools
[Workload Driver](#) and [MongoDB Multithreaded Performance Test Tool](#) were used to generate MongoDB workloads (insert/delete/update/find)
- Data processing: 1000 data points collected in more than 1 hour, removing zeros
- epsilon: 0.30
- minPts: 10
- adjusted_rand_score: 0.84



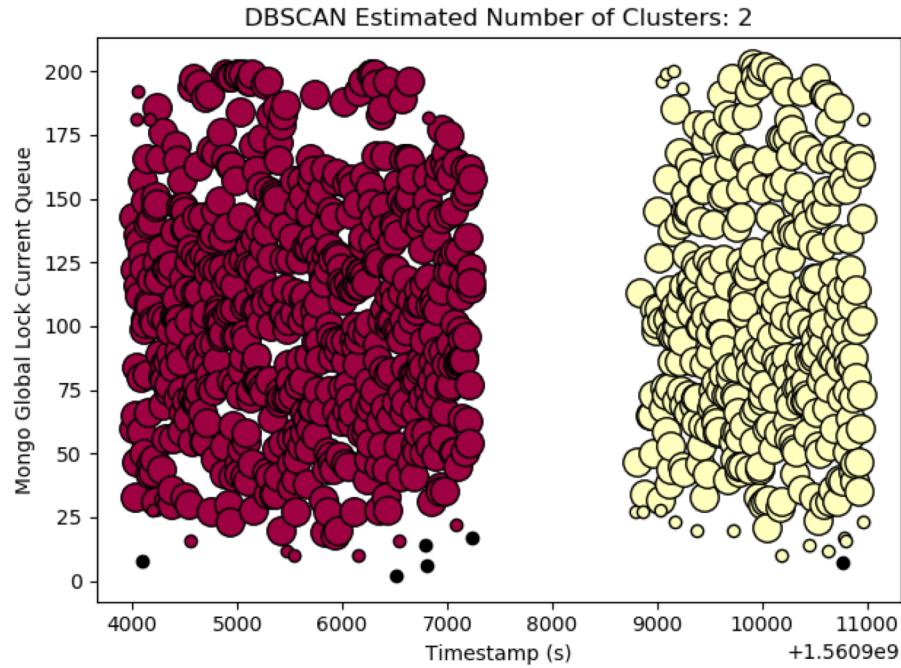
DBSCAN Graph - Node Exporter

- Node Exporter
 - Node Memory Available (MB)
- Data generation: MongoDB tools
[Workload Driver](#) and [MongoDB Multithreaded Performance Test Tool](#) were used to generate MongoDB workloads (insert/delete/update/find)
- Data processing: 1000 data points collected in more than 1 hour, removing zeros
- epsilon: 0.50
- minPts: 10
- adjusted_rand_score: 0.84



DBSCAN Graph - MongoDB

- MongoDB Global Lock Current Queue
- Data generation: MongoDB tools
[Workload Driver](#) and [MongoDB Multithreaded Performance Test Tool](#)
were used to generate MongoDB workloads (insert/delete/update/find)
- Data processing: 1000 data points collected in more than 1 hour, removing zeros
 - Gap between 2 clusters indicates a period with no data generated
- epsilon: 0.30
- minPts: 10
- adjusted_rand_score: 0.64



admin



default_region

Home

Resource statistics

Resource

Volumes / Buckets / File Share

Dataflow

Through migration / replication capability.

Monitor

Telemetry information.

Services

Orchestration services.

Profile

Profiles

Infrastructure

Regions, availability zones and storage

Identity

Managing tenants and users



Resource

1



Volumes

0



Buckets

0



Filesystems

Dataflow Quantity

0



Migrations

0



Replications

What's Next

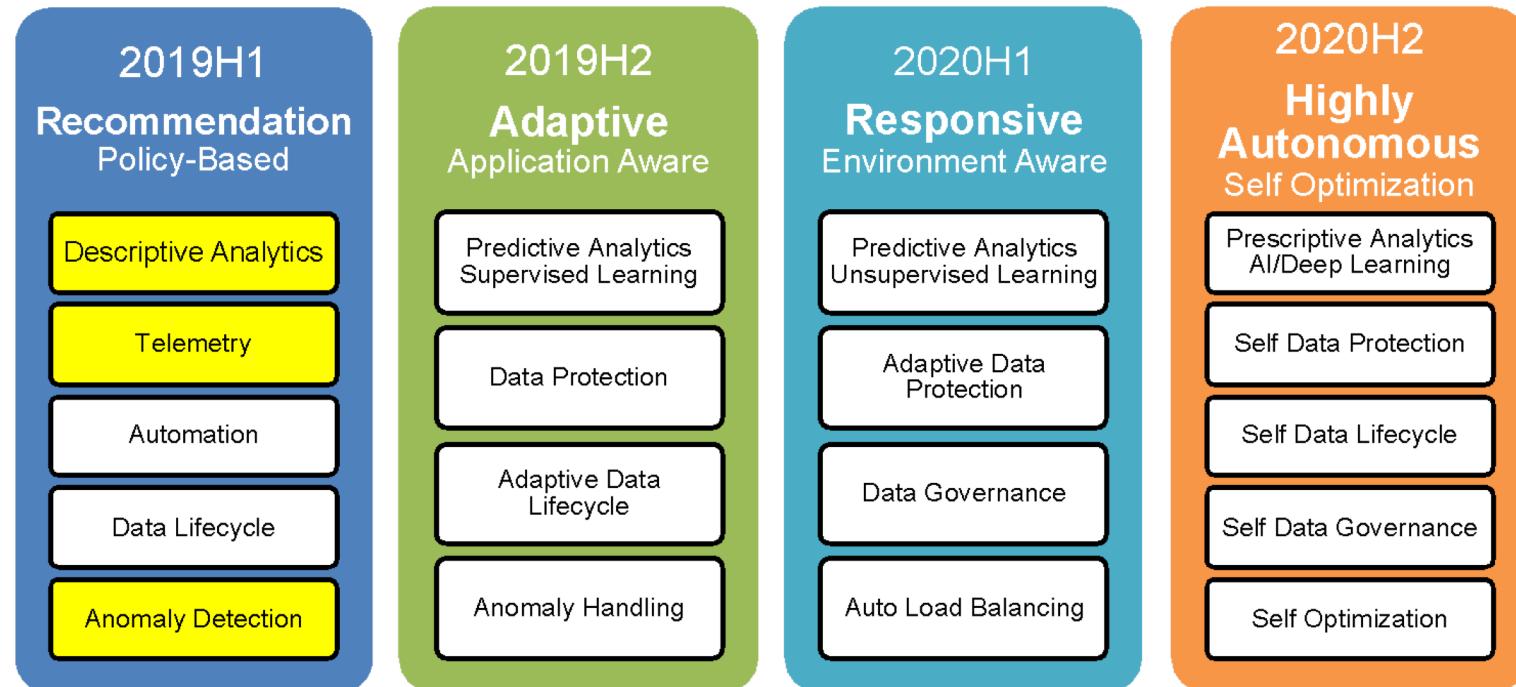
- Collect more data
- Correlate storage metrics, node-exporter metrics, with performance issues in applications running on storage provisioned by OpenSDS in Kubernetes environment
- Other algorithms to consider
 - Random forest
 - ARIMA - AutoRegressive Integrated Moving Average
 -
- Continue the journey towards self-driving storage ...



The Road To Autonomous Data Platform



China 2019



THANK YOU:



- <https://www.opensds.io>
- <https://github.com/opensds>
- info@opensds.io
- [@opensds_io](https://twitter.com/opensds_io)
- [opensds.slack.com](https://slack.com)

FIND OUT
MORE

JOIN AS
MEMBER

www.opensds.io

Accepting New Members
Vendors And End Users Welcome





KubeCon



CloudNativeCon

OPEN SOURCE SUMMIT

China 2019